

PA1_CS256_SP24_Ke_Xu

April 15, 2024

1 CSE 256: NLP UCSD PA1:

1.1 Text Classification with Logistic Regression and FF Networks(100points).

The goal of this assignment is to get experience developing text classifiers with with linear models and simple feedforward neural networks. You will see the standard pipeline used in many NLP tasks (reading in data, preprocessing, training, and testing).

- Part 1: PyTorch Basics (25 points)
- Part 2: Logistic Regression and Feedforward Neural Networks (60 points)
- Part 3: Exploration (20 points)

Data. You will using a dataset of movie review snippets taken from IMDB.

1.1.1 Due: April 22, 2024 at 10pm

IMPORTANT: After copying this notebook to your Google Drive, paste a link to it below. To get a publicly-accessible link, click the *Share* button at the top right, then click “Get shareable link” and copy the link.

Link: https://drive.google.com/file/d/16gQXMM2UrQsQibiMZ_mPIrazklsK2WDW/view?usp=sl

Notes:

Make sure to save the notebook as you go along.

Submission instructions are located at the bottom of the notebook.

The code should run fairly quickly (a couple of minutes at most even without a GPU), if it takes much longer than that, its likely that you have introduced an error.

1.2 Mount your Google Drive to Colab

Note: TODO: you need to specify your working foldername in this cell below:

```
[1]: # This mounts your Google Drive to the Colab VM.
from google.colab import drive
drive.mount('/content/drive')
```

```

# TODO: Enter the foldername in your Drive where you have saved the unzipped
# assignment folder, e.g. 'cse256/assignments/PA1/'
FOLDERNAME = None
FOLDERNAME = 'CSE256PAs/PA1'
assert FOLDERNAME is not None, "[!] Enter the foldername."

# Now that we've mounted your Drive, this ensures that
# the Python interpreter of the Colab VM can load
# python files from within it.
import sys
sys.path.append('/content/drive/My Drive/{}'.format(FOLDERNAME))

# This is later used to use the IMDB reviews
%cd /content/drive/My\ Drive/$FOLDERNAME/

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.
/content/drive/My Drive/CSE256PAs/PA1

2 Part 1: PyTorch Basics (25 Points)

We will use PyTorch, a machine learning framework, for the programming assignments in this course. The first part of this assignment focuses on PyTorch and how it is used for NLP. If you are new to [PyTorch](#), it is highly recommended to go to work through [the 60 minute tutorial](#)

##Question 1.1 (2.5 points)

In state-of-the-art NLP, words are represented by low-dimensional vectors, referred to as *embeddings*. When processing sequences such as sentences, movie, reviews, or entire paragraphs, word embeddings are used to compute a vector representation of the sequence, denoted by x . In the cell below, the embeddings for the words in the sequence “Alice talked to” are provided. Your task is to combine these embeddings into a single vector representation x , using [element-wise vector addition](#). This method is a simple way to obtain a sequence representation, namely, it is a *continuous bag-of-words (BoW) representation* of a sequence.

```

[2]: import torch
torch.set_printoptions(sci_mode=False)
# Seed the random number generator for reproducibility
torch.manual_seed(0)

input_sequence = 'I like NLP'

# Initialize an embedding matrix
# We have a vocabulary of 5 words, each represented by a 10-dimensional
  ↳ embedding vector.
embeddings = torch.nn.Embedding(num_embeddings=5, embedding_dim=10)
vocab = {'I': 0, 'like': 1, 'NLP': 2, 'classifiers': 3, '.': 4}

```

```

# Convert the word to integer indices. These indices will be used to
# retrieve the corresponding embeddings from the embedding matrix.
# In PyTorch, operations are performed on Tensor objects, so we need to convert
# the list of indices to a LongTensor.
indices = torch.LongTensor([vocab[w] for w in input_sequence.split()])
input_sequence_embs = embeddings(indices)
print('sequence embedding tensor size: ', input_sequence_embs.size())

# The input_sequence_embs tensor contains the embeddings for each word in the
# input sequence.
# The next step is to aggregate these embeddings into a single vector
# representation.
# You will use element-wise addition to do this.
# Write the code to add the embeddings element-wise and store the result in the
# variable "x".

# print(input_sequence_embs)
### YOUR CODE HERE!
# Replace with the actual computation
x = torch.sum(input_sequence_embs, dim=0)

### DO NOT MODIFY THE LINE BELOW
print('input sequence embedding sum (continuous BoW): ', x)

```

```

sequence embedding tensor size:  torch.Size([3, 10])
input sequence embedding sum (continuous BoW):  tensor([-0.1770, -2.3993,
-0.4721,  2.6568,  2.7157, -0.1408, -1.8421, -3.6277,
          2.2783,  1.1165], grad_fn=<SumBackward1>)

```

##Question 1.2 (2.5 points) Element-wise addition is not the best way to aggregate individual word embeddings in a sequence into a single vector representation (a process known as *composition*). State one significant limitation of using element-wise addition as a composition function for word embeddings? —

Write your answer here (2-3 sentences) One significant limitation is the loss of word order information. In natural language, the order of words can dramatically change the meaning of a sentence but element-wise addition treats all words equally, aggregating their embeddings without preserving any information about their sequence. This means that sentences with the same words in different orders would result in the same vector representation, potentially leading to ambiguities and inaccuracies in tasks that depend on understanding the precise meaning conveyed by the word order.

##Question 1.3 (5 points) The [softmax function](#) is used in nearly all the neural network architectures we will look at in this course. The softmax is computed on an n -dimensional vector $\langle x_1, x_2, \dots, x_n \rangle$ as $\text{softmax}(x_i) = \frac{e^{x_i}}{\sum_{1 \leq j \leq n} e^{x_j}}$. Given the sequence representation x we just computed, we can use the softmax function in combination with a linear projection using a matrix W to transform x into a probability distribution p over the next word, expressed as $p = \text{softmax}(Wx)$. Let's look at this in the cell below:

```
[3]: # Initialize a random matrix W of size 10x5. This will serve as the weight
      ↪ matrix
      # for the linear projection of the vector x into a 5-dimensional space.
      W = torch.rand(10, 5)

      # Project the vector x to a 5-dimensional space using the matrix W. This
      ↪ projection is achieved through
      # matrix multiplication. After the projection, apply the softmax function to
      ↪ the result,
      # which converts the 5-dimensional projected vector into a probability
      ↪ distribution.
      # You can find the softmax function in PyTorch's API (torch.nn.functional.
      ↪ softmax).
      # Store the resulting probability distribution in the variable "probs".

      ### YOUR CODE HERE
      # Replace with the actual computation
      import torch.nn.functional as F
      probs = F.softmax(torch.matmul(W.transpose(0, 1), x), dim=0)

      ### DO NOT MODIFY THE BELOW LINE!
      print('probability distribution', probs)
```

```
probability distribution tensor([0.0718, 0.0998, 0.1331, 0.6762, 0.0191],
grad_fn=<SoftmaxBackward0>)
```

##Question 1.4 (5 points)

In the example so far, we focused on a single sequence (“I like NLP”). However, in practical applications, it’s common to process multiple sequences simultaneously. This practice, known as *batching*, allows for more efficient use of GPU parallelism. In batching, each sequence is considered an example within a larger batch

For this question, you will perform redo the previous computation, but with a batch of two sequences instead of just one. The final output of this cell should be a 2x5 matrix, where each row represents a probability distribution for a sequence. **Important: Avoid using loops in your solution, as you will lose points.** The code should be fully vectorized.

```
[4]: import torch
      import torch.nn.functional as F

      # For this example, we replicate our previous sequence indices to create a
      ↪ simple batch.
      # Normally, each example in the batch would be different.
      batch_indices = torch.cat(2 * [indices]).reshape((2, 3))
      batch_embs = embeddings(batch_indices)
      print('Batch embedding tensor size: ', batch_embs.size())

      # To process the batch, follow these steps:
```

```

# Step 1: Aggregate the embeddings for each example in the batch into a single
↳ representation.
# This is done through element-wise addition. Use torch.sum with the
↳ appropriate 'dim' argument
# to sum across the sequence length (not the batch dimension).

# Step 2: Project each aggregated representation into a 5-dimensional space
↳ using the matrix W.
# This involves matrix multiplication, ensuring the resulting batch has the
↳ shape 2x5.

# Step 3: Apply the softmax function to the projected representations to obtain
↳ probability distributions.
# Each row in the output matrix should sum to 1, representing a probability
↳ distribution for each batch example.

### YOUR CODE HERE
# Replace with the actual computation
batch_aggregated = torch.sum(batch_embs, dim=1)
batch_projected = torch.matmul(batch_aggregated, W)
batch_probs = F.softmax(batch_projected, dim=1)

### DO NOT MODIFY THE BELOW LINE
print("Batch probability distributions:", batch_probs)

```

```

Batch embedding tensor size: torch.Size([2, 3, 10])
Batch probability distributions: tensor([[0.0718, 0.0998, 0.1331, 0.6762,
0.0191],
          [0.0718, 0.0998, 0.1331, 0.6762, 0.0191]], grad_fn=<SoftmaxBackward0>)

```

##Question 1.5 (5 points)

When processing a text sequence, how should the system handle words that are not present in the existing vocabulary? In the current implementation, the presence of such out-of-vocabulary words causes the code to fail, as in the cell below. To address this issue, a simple solution is to use the special token <UNK>, added to the vocabulary to serve as a placeholder for any unknown words.

Modify the indexing function to ensure that it checks each word against the known vocabulary and substitutes any out-of-vocabulary words with the <UNK> token. Make sure not to add any new words to the vocabulary except for the <UNK> token. Don't forget to adjust the embedding table.

```

[5]: import torch

torch.set_printoptions(sci_mode=False)
# Seed the random number generator for reproducibility
torch.manual_seed(0)

input_sequence = 'I like linear'

```

```

# Initialize an embedding matrix
# We have a vocabulary of 6 words, each represented by a 10-dimensional
  ↳ embedding vector.
embeddings = torch.nn.Embedding(num_embeddings=6, embedding_dim=10)
vocab = {'I': 0, 'like': 1, 'NLP': 2, 'classifiers': 3, '.': 4, '<UNK>': 5}

# indices = torch.LongTensor([vocab[w] for w in input_sequence.split()]) ###
  ↳ MODIFY THIS INDEXING
indices = torch.LongTensor([vocab[w] if w in vocab else vocab['<UNK>'] for w in
  ↳ input_sequence.split()])
input_sequence_embs = embeddings(indices)
print('sequence embedding tensor size: ', input_sequence_embs.size())

```

```

sequence embedding tensor size:  torch.Size([3, 10])

```

3 Part 2: Logistic Regression and Feedforward Neural Networks (60 points)

In this part, you are going to experiment with Logistic Regression and Feedforward Neural Networks. Run the starter code to train a two-layer fully connected neural network on the IMDB Sentiment Classification Dataset. The code provided below generates two plots that display the train accuracy and test accuracy. You will build on code to produce different variants.

```

[6]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import os
import time
import scipy.stats
import copy
import torch
from torch import nn
import torch.nn.functional as F
from torch.utils.data import Dataset
from torch.utils.data import DataLoader
from sklearn.feature_extraction.text import CountVectorizer

##### Neural network class
#
# Network of two fully connected layers
# with ReLU activation function and Softmax output
#####

class NN2(nn.Module):
    def __init__(self, input_size, hidden_size):

```

```

        super().__init__()
        self.fc1 = nn.Linear(input_size, hidden_size) # First fully connected
        layer.
        self.fc2 = nn.Linear(hidden_size, 2) # Second fully connected layer,
        outputting two classes.

        # Define the forward pass of the neural network.
        # x: The input tensor.
        def forward(self, x):
            x = F.relu(self.fc1(x)) # Apply ReLU activation function after the
        first layer.
            x = self.fc2(x) # Pass the result to the second layer.
            x = F.softmax(x, dim=1) # Apply Softmax to obtain output probabilities.
            return x

```

```

[7]: ##### ReviewsDataset class
#
# create a dataset to be used for training and evaluation
#####
# Function to read reviews from a directory
def read_reviews(directory, num_reviews=1000):
    reviews = []
    for filename in os.listdir(directory)[:num_reviews]: # Limit the number of
    files read
        with open(os.path.join(directory, filename), 'r', encoding='utf-8') as
    file:
        reviews.append(file.read())
    return reviews

class ReviewsDataset(Dataset):
    def __init__(self, pos_dir, neg_dir, num_reviews=1000, vectorizer=None,
    train=True):
        self.reviews = read_reviews(pos_dir, num_reviews) +
    read_reviews(neg_dir, num_reviews)
        self.labels = [1] * min(num_reviews, len(os.listdir(pos_dir))) + [0] *
    min(num_reviews, len(os.listdir(neg_dir)))
        if train or vectorizer is None:
            self.vectorizer = CountVectorizer(max_features=512) # Adjust as
    needed
            self.embeddings = self.vectorizer.fit_transform(self.reviews).
    toarray()
        else:
            self.vectorizer = vectorizer
            self.embeddings = self.vectorizer.transform(self.reviews).toarray()

```

```

def __len__(self):
    return len(self.reviews)

def __getitem__(self, idx):
    return self.embeddings[idx], self.labels[idx]

```

```

[8]: ##### train_epoch
#
# function that trains for one epoch (one pass through the training set)
#####

def train_epoch(data_loader, model, loss_fn, optimizer):
    size = len(data_loader.dataset)
    num_batches = len(data_loader)
    model.train()
    train_loss, correct = 0, 0
    for batch, (X, y) in enumerate(data_loader):
        X = X.float()
        # Compute prediction error
        pred = model(X)
        loss = loss_fn(pred, y)
        train_loss += loss.item()
        correct += (pred.argmax(1) == y).type(torch.float).sum().item()

        # Backpropagation
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    average_train_loss = train_loss / num_batches
    accuracy = correct / size
    return accuracy, average_train_loss

##### eval_epoch
#
# function that evaluates a model with a test set
#####

def eval_epoch(data_loader, model, loss_fn, optimizer):
    size = len(data_loader.dataset)
    num_batches = len(data_loader)
    model.eval()
    eval_loss = 0
    correct = 0
    for batch, (X, y) in enumerate(data_loader):
        # Compute prediction error
        X = X.float()

```



```

    pred = model(X)
    loss = loss_fn(pred, y)
    eval_loss += loss.item()
    correct += (pred.argmax(1) == y).type(torch.float).sum().item()

    average_eval_loss = eval_loss / num_batches
    accuracy = correct / size
    return accuracy, average_eval_loss

##### experiment
#
# function that trains a neural network with a training set
# and evaluates the neural network with a test set
#####
def experiment(model):

    # negative log likelihood loss function
    loss_fn = nn.NLLLoss()

    # Adam optimizer
    optimizer = torch.optim.Adam(model.parameters(), lr=0.0001)

    average_train_loss = []
    all_train_accuracy = []
    average_test_loss = []
    all_test_accuracy = []
    for epoch in range(150):
        train_accuracy, train_loss = train_epoch(train_loader, model,
↳loss_fn, optimizer)
        all_train_accuracy += [train_accuracy]
        test_accuracy, test_loss = eval_epoch(test_loader, model,
↳loss_fn, optimizer)
        all_test_accuracy += [test_accuracy]
        if epoch % 10 == 9:
            print(f'Epoch #{epoch+1}: \t train accuracy_
↳{train_accuracy:.3f}\t train loss {train_loss:.3f}\t test accuracy_
↳{test_accuracy:.3f}\t test loss {test_loss:.3f}')
    return all_train_accuracy, all_test_accuracy

```

```

[9]: #####
#
# 1) Load data splits: the train and test sets
# 2) Train neural networks
# 3) Plot the results
#####

```

```

start_time = time.time()

# Load the dataset
root_dir = '/content/drive/My Drive/{}/aclImdb/train'.format(FOLDERNAME)
root_dir_test = '/content/drive/My Drive/{}/aclImdb/test'.format(FOLDERNAME)
train_dataset = ReviewsDataset(root_dir+'/pos', root_dir+'/neg', train=True)
test_dataset = ReviewsDataset(root_dir_test+'/pos', root_dir_test+'/neg',
    ↪vectorizer=train_dataset.vectorizer, train=False)

train_loader = DataLoader(train_dataset, batch_size=16, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=16, shuffle=False)

end_time = time.time()
elapsed_time = end_time - start_time

print(f"Time to load data: {elapsed_time} seconds")

```

Time to load data: 166.29224681854248 seconds

```

[10]: start_time = time.time()

# train neural networks
print('\n2 layers:')
nn2_train_accuracy, nn2_test_accuracy = experiment(NN2(input_size=512,
    ↪hidden_size=100))

# plot training accuracy
plt.plot(nn2_train_accuracy)
plt.title('training accuracy (varying # of layers)')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['2 layers'])
plt.show()

# plot testing accuracy
plt.plot(nn2_test_accuracy)
plt.title('testing accuracy (varying # of layers)')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['2 layers'])
plt.show()


end_time = time.time()
elapsed_time = end_time - start_time


print(f"Time to train, eval model: {elapsed_time} seconds")

```

2 layers:

Epoch #10:	train accuracy 0.881	train loss -0.821	test accuracy
0.793	test loss -0.745		
Epoch #20:	train accuracy 0.924	train loss -0.887	test accuracy
0.790	test loss -0.767		
Epoch #30:	train accuracy 0.947	train loss -0.917	test accuracy
0.786	test loss -0.773		
Epoch #40:	train accuracy 0.955	train loss -0.935	test accuracy
0.783	test loss -0.774		
Epoch #50:	train accuracy 0.961	train loss -0.947	test accuracy
0.783	test loss -0.776		
Epoch #60:	train accuracy 0.964	train loss -0.954	test accuracy
0.782	test loss -0.776		
Epoch #70:	train accuracy 0.967	train loss -0.961	test accuracy
0.778	test loss -0.775		
Epoch #80:	train accuracy 0.969	train loss -0.965	test accuracy
0.780	test loss -0.776		
Epoch #90:	train accuracy 0.970	train loss -0.967	test accuracy
0.778	test loss -0.777		
Epoch #100:	train accuracy 0.970	train loss -0.968	test accuracy
0.775	test loss -0.776		
Epoch #110:	train accuracy 0.970	train loss -0.969	test accuracy
0.774	test loss -0.776		
Epoch #120:	train accuracy 0.970	train loss -0.969	test accuracy
0.776	test loss -0.776		
Epoch #130:	train accuracy 0.970	train loss -0.969	test accuracy
0.773	test loss -0.776		
Epoch #140:	train accuracy 0.970	train loss -0.969	test accuracy
0.773	test loss -0.776		
Epoch #150:	train accuracy 0.970	train loss -0.969	test accuracy
0.774	test loss -0.776		





Time to train, eval model: 46.318968296051025 seconds

3.0.1 TO DO: Impelementation

- Implement and test fully connected networks with 1,2,3, and 4 layers. The starter code above already provides you with an implementation of 2 layers. Each hidden layer should have 100 nodes.
- On the four layer network, modify the code to replace the ReLU activation function with the sigmoid activation function.
- On the four layer network, modify your code to insert a dropout layer with probability 0.5 after each hidden leayer. Tip: see the function `nn.dropout()`.

```
[11]: ### YOUR CODE HERE
import torch
from torch import nn
import torch.nn.functional as F

class NN1(nn.Module):
    def __init__(self, input_size):
        super(NN1, self).__init__()
```

```

        self.fc1 = nn.Linear(input_size, 2)  # Single layer directly to output

    def forward(self, x):
        x = F.log_softmax(self.fc1(x), dim=1)
        return x

class NN2(nn.Module):
    def __init__(self, input_size, hidden_size=100):
        super(NN2, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.fc2 = nn.Linear(hidden_size, 2)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.log_softmax(self.fc2(x), dim=1)
        return x

class NN3(nn.Module):
    def __init__(self, input_size, hidden_size=100):
        super(NN3, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.fc2 = nn.Linear(hidden_size, hidden_size)
        self.fc3 = nn.Linear(hidden_size, 2)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.log_softmax(self.fc3(x), dim=1)
        return x

class NN4(nn.Module):
    def __init__(self, input_size, hidden_size=100):
        super(NN4, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.fc2 = nn.Linear(hidden_size, hidden_size)
        self.fc3 = nn.Linear(hidden_size, hidden_size)
        self.fc4 = nn.Linear(hidden_size, 2)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = F.relu(self.fc3(x))
        x = F.log_softmax(self.fc4(x), dim=1)
        return x

input_size = 512
hidden_size = 100

```

```

models = [NN1(input_size), NN2(input_size, hidden_size), NN3(input_size,
↳hidden_size), NN4(input_size, hidden_size)]
model_names = ['1 layer', '2 layers', '3 layers', '4 layers']

# Dictionaries to hold training and testing accuracies
train_accuracies = {}
test_accuracies = {}

start_time = time.time()

# Loop through each model, train, and collect accuracies
for model, name in zip(models, model_names):
    print(f'\nTraining {name}:')
    train_accuracy, test_accuracy = experiment(model)

    train_accuracies[name] = train_accuracy
    test_accuracies[name] = test_accuracy

end_time = time.time()
elapsed_time = end_time - start_time
print(f"\nTime to train and evaluate all models: {elapsed_time} seconds")

# Plot training accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(train_accuracies[name], label=f'{name}')
plt.title('Training Accuracy by Number of Layers')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy')
plt.legend()
plt.show()

# Plot testing accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(test_accuracies[name], label=f'{name}')
plt.title('Testing Accuracy by Number of Layers')
plt.xlabel('Epoch')
plt.ylabel('Testing Accuracy')
plt.legend()
plt.show()

```

Training 1 layer:

Epoch #10:	train accuracy 0.796	train loss 0.497	test accuracy
0.757	test loss 0.548		
Epoch #20:	train accuracy 0.848	train loss 0.412	test accuracy
0.781	test loss 0.504		

Epoch #30:	train accuracy 0.875	train loss 0.363	test accuracy
0.785	test loss 0.482		
Epoch #40:	train accuracy 0.883	train loss 0.331	test accuracy
0.789	test loss 0.472		
Epoch #50:	train accuracy 0.887	train loss 0.309	test accuracy
0.788	test loss 0.469		
Epoch #60:	train accuracy 0.899	train loss 0.292	test accuracy
0.790	test loss 0.465		
Epoch #70:	train accuracy 0.900	train loss 0.278	test accuracy
0.787	test loss 0.469		
Epoch #80:	train accuracy 0.904	train loss 0.267	test accuracy
0.790	test loss 0.470		
Epoch #90:	train accuracy 0.911	train loss 0.258	test accuracy
0.788	test loss 0.475		
Epoch #100:	train accuracy 0.910	train loss 0.250	test accuracy
0.785	test loss 0.483		
Epoch #110:	train accuracy 0.913	train loss 0.244	test accuracy
0.782	test loss 0.485		
Epoch #120:	train accuracy 0.913	train loss 0.238	test accuracy
0.779	test loss 0.487		
Epoch #130:	train accuracy 0.916	train loss 0.233	test accuracy
0.779	test loss 0.501		
Epoch #140:	train accuracy 0.918	train loss 0.228	test accuracy
0.777	test loss 0.503		
Epoch #150:	train accuracy 0.921	train loss 0.223	test accuracy
0.777	test loss 0.505		

Training 2 layers:

Epoch #10:	train accuracy 0.893	train loss 0.296	test accuracy
0.797	test loss 0.471		
Epoch #20:	train accuracy 0.933	train loss 0.195	test accuracy
0.784	test loss 0.516		
Epoch #30:	train accuracy 0.962	train loss 0.135	test accuracy
0.778	test loss 0.581		
Epoch #40:	train accuracy 0.983	train loss 0.090	test accuracy
0.773	test loss 0.653		
Epoch #50:	train accuracy 0.995	train loss 0.058	test accuracy
0.768	test loss 0.756		
Epoch #60:	train accuracy 0.998	train loss 0.036	test accuracy
0.765	test loss 0.863		
Epoch #70:	train accuracy 0.999	train loss 0.021	test accuracy
0.765	test loss 0.953		
Epoch #80:	train accuracy 1.000	train loss 0.012	test accuracy
0.764	test loss 1.078		
Epoch #90:	train accuracy 1.000	train loss 0.007	test accuracy
0.766	test loss 1.182		
Epoch #100:	train accuracy 1.000	train loss 0.004	test accuracy
0.768	test loss 1.312		

Epoch #110:	train accuracy 1.000	train loss 0.002	test accuracy
0.769	test loss 1.413		
Epoch #120:	train accuracy 1.000	train loss 0.001	test accuracy
0.761	test loss 1.509		
Epoch #130:	train accuracy 1.000	train loss 0.001	test accuracy
0.764	test loss 1.580		
Epoch #140:	train accuracy 1.000	train loss 0.001	test accuracy
0.764	test loss 1.653		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.765	test loss 1.737		

Training 3 layers:


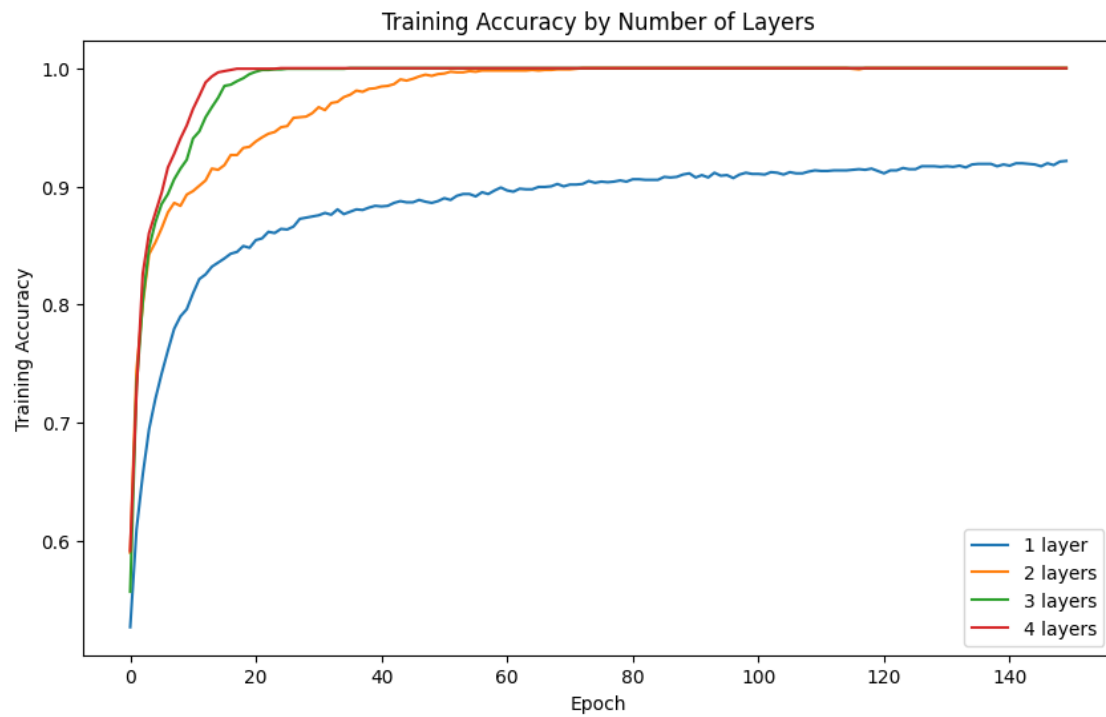
Epoch #10:	train accuracy 0.922	train loss 0.211	test accuracy
0.792	test loss 0.510		
Epoch #20:	train accuracy 0.995	train loss 0.045	test accuracy
0.781	test loss 0.720		
Epoch #30:	train accuracy 1.000	train loss 0.009	test accuracy
0.778	test loss 0.961		
Epoch #40:	train accuracy 1.000	train loss 0.002	test accuracy
0.778	test loss 1.163		
Epoch #50:	train accuracy 1.000	train loss 0.001	test accuracy
0.775	test loss 1.348		
Epoch #60:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 1.495		
Epoch #70:	train accuracy 1.000	train loss 0.000	test accuracy
0.772	test loss 1.660		
Epoch #80:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 1.784		
Epoch #90:	train accuracy 1.000	train loss 0.000	test accuracy
0.771	test loss 1.911		
Epoch #100:	train accuracy 1.000	train loss 0.000	test accuracy
0.771	test loss 2.039		
Epoch #110:	train accuracy 1.000	train loss 0.000	test accuracy
0.769	test loss 2.159		
Epoch #120:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 2.279		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.767	test loss 2.406		
Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 2.527		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.767	test loss 2.648		

Training 4 layers:

Epoch #10:	train accuracy 0.952	train loss 0.141	test accuracy
0.783	test loss 0.601		
Epoch #20:	train accuracy 1.000	train loss 0.009	test accuracy
0.780	test loss 1.098		

Epoch #30:	train accuracy 1.000	train loss 0.001	test accuracy
0.774	test loss 1.442		
Epoch #40:	train accuracy 1.000	train loss 0.000	test accuracy
0.777	test loss 1.662		
Epoch #50:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 1.867		
Epoch #60:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 2.048		
Epoch #70:	train accuracy 1.000	train loss 0.000	test accuracy
0.776	test loss 2.214		
Epoch #80:	train accuracy 1.000	train loss 0.000	test accuracy
0.777	test loss 2.374		
Epoch #90:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 2.532		
Epoch #100:	train accuracy 1.000	train loss 0.000	test accuracy
0.776	test loss 2.676		
Epoch #110:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.830		
Epoch #120:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.962		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 3.097		
Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 3.237		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 3.365		

Time to train and evaluate all models: 180.06459021568298 seconds



```
[12]: import torch
from torch import nn
import torch.nn.functional as F

class NN4_Sigmoid_Dropout(nn.Module):
    def __init__(self, input_size, hidden_size=100):
        super(NN4_Sigmoid_Dropout, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.dropout1 = nn.Dropout(0.5)
        self.fc2 = nn.Linear(hidden_size, hidden_size)
        self.dropout2 = nn.Dropout(0.5)
        self.fc3 = nn.Linear(hidden_size, hidden_size)
        self.dropout3 = nn.Dropout(0.5)
        self.fc4 = nn.Linear(hidden_size, 2)

    def forward(self, x):
        x = torch.sigmoid(self.fc1(x))
        x = self.dropout1(x)
        x = torch.sigmoid(self.fc2(x))
        x = self.dropout2(x)
        x = torch.sigmoid(self.fc3(x))
        x = self.dropout3(x)
        x = F.log_softmax(self.fc4(x), dim=1)
        return x
```

3.1 Question 2.1 Architecture Comparison (20 points)

Generate two plots where the y-axis is the accuracy and the x-axis is the # of epochs. The first plot should include 4 curves that show the training accuracy for 1, 2, 3, and 4 layers. The second plot should include 4 curves that show the testing accuracy for 1, 2, 3, and 4 layers. Use ReLU activation functions without any dropout and 100 nodes per hidden layer. Discuss the results.

```
[13]: ### YOUR CODE HERE
input_size = 512
hidden_size = 100
models = [NN1(input_size), NN2(input_size, hidden_size), NN3(input_size,
↪hidden_size), NN4(input_size, hidden_size), NN4_Sigmoid_Dropout(input_size,
↪hidden_size)]
model_names = ['1 layer', '2 layers', '3 layers', '4 layers', '4 layers Sigmoid
↪with Dropout 0.5']

# Dictionaries to hold training and testing accuracies
train_accuracies = {}
test_accuracies = {}

start_time = time.time()
```

```

# Loop through each model, train, and collect accuracies
for model, name in zip(models, model_names):
    print(f'\nTraining {name}:')
    train_accuracy, test_accuracy = experiment(model)

    train_accuracies[name] = train_accuracy
    test_accuracies[name] = test_accuracy

end_time = time.time()
elapsed_time = end_time - start_time
print(f"\nTime to train and evaluate all models: {elapsed_time} seconds")

# Plot training accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(train_accuracies[name], label=f'{name}')
plt.title('Training Accuracy by Number of Layers')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy')
plt.legend()
plt.show()

# Plot testing accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(test_accuracies[name], label=f'{name}')
plt.title('Testing Accuracy by Number of Layers')
plt.xlabel('Epoch')
plt.ylabel('Testing Accuracy')
plt.legend()
plt.show()

```

Training 1 layer:

Epoch #10:	train accuracy 0.811	train loss 0.496	test accuracy
0.734	test loss 0.565		
Epoch #20:	train accuracy 0.859	train loss 0.408	test accuracy
0.781	test loss 0.513		
Epoch #30:	train accuracy 0.875	train loss 0.360	test accuracy
0.791	test loss 0.489		
Epoch #40:	train accuracy 0.885	train loss 0.329	test accuracy
0.794	test loss 0.476		
Epoch #50:	train accuracy 0.893	train loss 0.307	test accuracy
0.793	test loss 0.471		
Epoch #60:	train accuracy 0.895	train loss 0.291	test accuracy
0.793	test loss 0.471		
Epoch #70:	train accuracy 0.900	train loss 0.278	test accuracy
0.791	test loss 0.470		

Epoch #80:	train accuracy 0.906	train loss 0.267	test accuracy
0.790	test loss 0.473		
Epoch #90:	train accuracy 0.906	train loss 0.258	test accuracy
0.792	test loss 0.478		
Epoch #100:	train accuracy 0.910	train loss 0.250	test accuracy
0.789	test loss 0.482		
Epoch #110:	train accuracy 0.912	train loss 0.243	test accuracy
0.790	test loss 0.485		
Epoch #120:	train accuracy 0.915	train loss 0.238	test accuracy
0.784	test loss 0.497		
Epoch #130:	train accuracy 0.915	train loss 0.232	test accuracy
0.781	test loss 0.497		
Epoch #140:	train accuracy 0.918	train loss 0.227	test accuracy
0.779	test loss 0.502		
Epoch #150:	train accuracy 0.922	train loss 0.223	test accuracy
0.778	test loss 0.506		

Training 2 layers:

Epoch #10:	train accuracy 0.891	train loss 0.295	test accuracy
0.787	test loss 0.475		
Epoch #20:	train accuracy 0.940	train loss 0.187	test accuracy
0.793	test loss 0.509		
Epoch #30:	train accuracy 0.974	train loss 0.120	test accuracy
0.781	test loss 0.572		
Epoch #40:	train accuracy 0.990	train loss 0.075	test accuracy
0.775	test loss 0.650		
Epoch #50:	train accuracy 0.997	train loss 0.045	test accuracy
0.775	test loss 0.763		
Epoch #60:	train accuracy 0.999	train loss 0.026	test accuracy
0.777	test loss 0.829		
Epoch #70:	train accuracy 1.000	train loss 0.015	test accuracy
0.773	test loss 0.950		
Epoch #80:	train accuracy 1.000	train loss 0.008	test accuracy
0.773	test loss 1.049		
Epoch #90:	train accuracy 1.000	train loss 0.005	test accuracy
0.770	test loss 1.158		
Epoch #100:	train accuracy 1.000	train loss 0.002	test accuracy
0.774	test loss 1.271		
Epoch #110:	train accuracy 1.000	train loss 0.001	test accuracy
0.774	test loss 1.392		
Epoch #120:	train accuracy 1.000	train loss 0.001	test accuracy
0.777	test loss 1.490		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.772	test loss 1.624		
Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 1.755		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 1.873		

Training 3 layers:

Epoch #10:	train accuracy 0.926	train loss 0.207	test accuracy
0.783	test loss 0.519		
Epoch #20:	train accuracy 0.993	train loss 0.051	test accuracy
0.778	test loss 0.750		
Epoch #30:	train accuracy 1.000	train loss 0.010	test accuracy
0.777	test loss 1.007		
Epoch #40:	train accuracy 1.000	train loss 0.002	test accuracy
0.775	test loss 1.217		
Epoch #50:	train accuracy 1.000	train loss 0.001	test accuracy
0.775	test loss 1.421		
Epoch #60:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 1.573		
Epoch #70:	train accuracy 1.000	train loss 0.000	test accuracy
0.776	test loss 1.728		
Epoch #80:	train accuracy 1.000	train loss 0.000	test accuracy
0.773	test loss 1.872		
Epoch #90:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.006		
Epoch #100:	train accuracy 1.000	train loss 0.000	test accuracy
0.776	test loss 2.161		
Epoch #110:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.268		
Epoch #120:	train accuracy 1.000	train loss 0.000	test accuracy
0.776	test loss 2.410		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 2.531		
Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.658		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.773	test loss 2.791		

Training 4 layers:

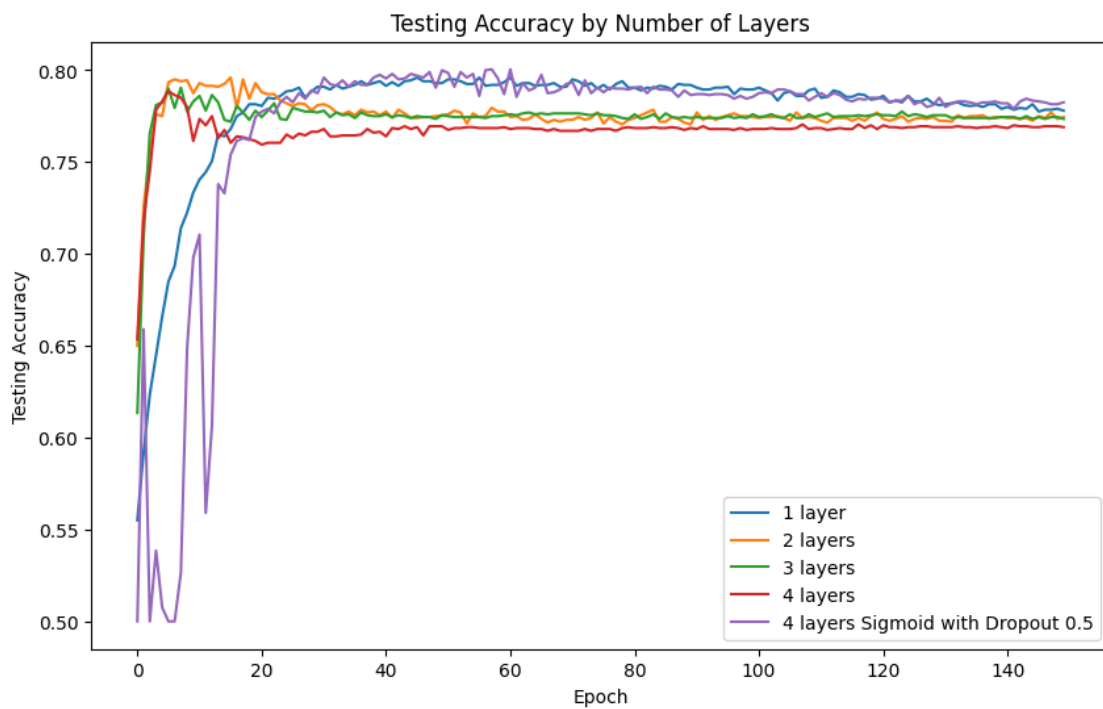

Epoch #10:	train accuracy 0.942	train loss 0.164	test accuracy
0.761	test loss 0.673		
Epoch #20:	train accuracy 0.999	train loss 0.012	test accuracy
0.761	test loss 1.077		
Epoch #30:	train accuracy 1.000	train loss 0.001	test accuracy
0.766	test loss 1.430		
Epoch #40:	train accuracy 1.000	train loss 0.000	test accuracy
0.766	test loss 1.713		
Epoch #50:	train accuracy 1.000	train loss 0.000	test accuracy
0.769	test loss 1.909		
Epoch #60:	train accuracy 1.000	train loss 0.000	test accuracy
0.769	test loss 2.082		
Epoch #70:	train accuracy 1.000	train loss 0.000	test accuracy
0.767	test loss 2.266		

Epoch #80:	train accuracy 1.000	train loss 0.000	test accuracy
0.769	test loss 2.447		
Epoch #90:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 2.606		
Epoch #100:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 2.759		
Epoch #110:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 2.914		
Epoch #120:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 3.049		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.769	test loss 3.195		
Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.769	test loss 3.331		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.769	test loss 3.456		

Training 4 layers Sigmoid with Dropout 0.5:

Epoch #10:	train accuracy 0.528	train loss 0.699	test accuracy
0.698	test loss 0.688		
Epoch #20:	train accuracy 0.734	train loss 0.579	test accuracy
0.774	test loss 0.576		
Epoch #30:	train accuracy 0.848	train loss 0.369	test accuracy
0.787	test loss 0.452		
Epoch #40:	train accuracy 0.856	train loss 0.343	test accuracy
0.797	test loss 0.449		
Epoch #50:	train accuracy 0.878	train loss 0.297	test accuracy
0.800	test loss 0.459		
Epoch #60:	train accuracy 0.882	train loss 0.289	test accuracy
0.787	test loss 0.475		
Epoch #70:	train accuracy 0.895	train loss 0.266	test accuracy
0.792	test loss 0.483		
Epoch #80:	train accuracy 0.889	train loss 0.263	test accuracy
0.789	test loss 0.496		
Epoch #90:	train accuracy 0.903	train loss 0.248	test accuracy
0.786	test loss 0.512		
Epoch #100:	train accuracy 0.905	train loss 0.236	test accuracy
0.785	test loss 0.517		
Epoch #110:	train accuracy 0.913	train loss 0.227	test accuracy
0.785	test loss 0.532		
Epoch #120:	train accuracy 0.924	train loss 0.212	test accuracy
0.785	test loss 0.549		
Epoch #130:	train accuracy 0.918	train loss 0.207	test accuracy
0.785	test loss 0.559		
Epoch #140:	train accuracy 0.912	train loss 0.222	test accuracy
0.782	test loss 0.580		
Epoch #150:	train accuracy 0.916	train loss 0.209	test accuracy
0.782	test loss 0.596		

Time to train and evaluate all models: 243.20409536361694 seconds



Analysis and discussion here (< 5 sentences) : The training accuracy graph indicates that models with more layers tend to overfit, as seen with the three-layer network achieving the highest training accuracy but then plateauing. The four-layer network with dropout regularisation doesn't achieve as high training accuracy as the others, likely due to its regularization effect preventing overfitting. In contrast, the testing accuracy graph shows that the models with two and three layers have similar performance and generalize well after a certain number of epochs. The four-layer network with dropout has significantly improved testing performance compared to the same network without dropout, indicating the effectiveness of dropout in combating overfitting.

3.2 Question 2.2 Activation functions (20 points)

Generate two plots where the y-axis is the accuracy and the x-axis is the # of epochs. The first plot should include 2 curves that show the training accuracy when using the ReLU versus sigmoid activation functions. The second plot should include 2 curves that show the testing accuracy when using the ReLU versus sigmoid activation functions. Use 2 layers and 100 nodes per hidden layer without any dropout. Discuss the results.

```
[14]: ### YOUR CODE HERE
class NN2_Sigmoid(nn.Module):
    def __init__(self, input_size, hidden_size=100):
        super(NN2_Sigmoid, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.fc2 = nn.Linear(hidden_size, 2)

    def forward(self, x):
        x = torch.sigmoid(self.fc1(x))
        x = F.log_softmax(self.fc2(x), dim=1)
        return x

input_size = 512
hidden_size = 100
models = [NN2(input_size), NN2_Sigmoid(input_size, hidden_size)]
model_names = ['2 layer RELU', '2 layers Sigmoid']

# Dictionaries to hold training and testing accuracies
train_accuracies = {}
test_accuracies = {}

start_time = time.time()

# Loop through each model, train, and collect accuracies
for model, name in zip(models, model_names):
    print(f'\nTraining {name}:')
    train_accuracy, test_accuracy = experiment(model)
```

```

train_accuracies[name] = train_accuracy
test_accuracies[name] = test_accuracy

end_time = time.time()
elapsed_time = end_time - start_time
print(f"\nTime to train and evaluate all models: {elapsed_time} seconds")

# Plot training accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(train_accuracies[name], label=f'{name}')
plt.title('Training Accuracy by Activation Methods')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy')
plt.legend()
plt.show()

# Plot testing accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(test_accuracies[name], label=f'{name}')
plt.title('Testing Accuracy by Activation Methods')
plt.xlabel('Epoch')
plt.ylabel('Testing Accuracy')
plt.legend()
plt.show()

```

Training 2 layer RELU:



Epoch #10:	train accuracy 0.894	train loss 0.289	test accuracy 0.789	test loss 0.473
Epoch #20:	train accuracy 0.944	train loss 0.184	test accuracy 0.787	test loss 0.503
Epoch #30:	train accuracy 0.972	train loss 0.119	test accuracy 0.779	test loss 0.592
Epoch #40:	train accuracy 0.989	train loss 0.076	test accuracy 0.779	test loss 0.668
Epoch #50:	train accuracy 0.998	train loss 0.046	test accuracy 0.771	test loss 0.775
Epoch #60:	train accuracy 0.999	train loss 0.027	test accuracy 0.773	test loss 0.861
Epoch #70:	train accuracy 1.000	train loss 0.015	test accuracy 0.768	test loss 0.990
Epoch #80:	train accuracy 1.000	train loss 0.008	test accuracy 0.770	test loss 1.074
Epoch #90:	train accuracy 1.000	train loss 0.005	test accuracy 0.773	test loss 1.193

Epoch #100:	train accuracy 1.000	train loss 0.002	test accuracy
0.771	test loss 1.320		
Epoch #110:	train accuracy 1.000	train loss 0.001	test accuracy
0.769	test loss 1.459		
Epoch #120:	train accuracy 1.000	train loss 0.001	test accuracy
0.766	test loss 1.603		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 1.724		
Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 1.827		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.768	test loss 1.963		

Training 2 layers Sigmoid:

Epoch #10:	train accuracy 0.861	train loss 0.424	test accuracy
0.786	test loss 0.501		
Epoch #20:	train accuracy 0.884	train loss 0.305	test accuracy
0.799	test loss 0.454		
Epoch #30:	train accuracy 0.903	train loss 0.257	test accuracy
0.794	test loss 0.460		
Epoch #40:	train accuracy 0.913	train loss 0.227	test accuracy
0.789	test loss 0.487		
Epoch #50:	train accuracy 0.923	train loss 0.206	test accuracy
0.782	test loss 0.510		
Epoch #60:	train accuracy 0.930	train loss 0.191	test accuracy
0.774	test loss 0.540		
Epoch #70:	train accuracy 0.935	train loss 0.177	test accuracy
0.773	test loss 0.576		
Epoch #80:	train accuracy 0.944	train loss 0.167	test accuracy
0.770	test loss 0.618		
Epoch #90:	train accuracy 0.945	train loss 0.158	test accuracy
0.770	test loss 0.635		
Epoch #100:	train accuracy 0.953	train loss 0.148	test accuracy
0.770	test loss 0.676		
Epoch #110:	train accuracy 0.957	train loss 0.141	test accuracy
0.764	test loss 0.703		
Epoch #120:	train accuracy 0.958	train loss 0.133	test accuracy
0.763	test loss 0.738		
Epoch #130:	train accuracy 0.964	train loss 0.127	test accuracy
0.765	test loss 0.770		
Epoch #140:	train accuracy 0.965	train loss 0.118	test accuracy
0.763	test loss 0.804		
Epoch #150:	train accuracy 0.970	train loss 0.112	test accuracy
0.761	test loss 0.840		

Time to train and evaluate all models: 84.39417600631714 seconds



Analysis and discussion here (< 5 sentences) : The training accuracy plot reveals that the network with ReLU activation converges faster and achieves marginally higher accuracy compared to the sigmoid network. In testing accuracy, both ReLU and sigmoid display similar patterns of convergence, with ReLU having a slight advantage. The ReLU's performance suggests it might be more effective at capturing complex patterns in this context. However, the close performance on the testing set suggests that both activation functions are capable of similar generalization from the training data, and the choice may depend on the specific characteristics of the dataset and the training dynamics.

3.3 Question 2.3 Dropout comparison (15 points)

Generate two plots where the y-axis is the accuracy and the x-axis is the # of epochs. The first plot should include 2 curves that show the training accuracy with and without dropout (with probability 0.5) after each hidden layer. The second plot should include 2 curves that show the testing accuracy with and without dropout (with probability 0.5) after each hidden layer. Use 4 layers and 36 nodes per hidden layer with ReLU activation functions. Discuss the results.

```
[15]: ### YOUR CODE HERE
class NN4_Dropout(nn.Module):
    def __init__(self, input_size, hidden_size=100):
        super(NN4_Dropout, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.dropout1 = nn.Dropout(0.5)
        self.fc2 = nn.Linear(hidden_size, hidden_size)
        self.dropout2 = nn.Dropout(0.5)
        self.fc3 = nn.Linear(hidden_size, hidden_size)
        self.dropout3 = nn.Dropout(0.5)
        self.fc4 = nn.Linear(hidden_size, 2)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.dropout1(x)
        x = F.relu(self.fc2(x))
        x = self.dropout2(x)
        x = F.relu(self.fc3(x))
        x = self.dropout3(x)
        x = F.log_softmax(self.fc4(x), dim=1)
        return x

input_size = 512
hidden_size = 36
models = [NN4(input_size, hidden_size), NN4_Dropout(input_size, hidden_size)]
model_names = ['4 layer RELU', '4 layers RELU with Dropout']

# Dictionaries to hold training and testing accuracies
train_accuracies = {}
test_accuracies = {}
```

```

start_time = time.time()

# Loop through each model, train, and collect accuracies
for model, name in zip(models, model_names):
    print(f'\nTraining {name}:')
    train_accuracy, test_accuracy = experiment(model)

    train_accuracies[name] = train_accuracy
    test_accuracies[name] = test_accuracy

end_time = time.time()
elapsed_time = end_time - start_time
print(f"\nTime to train and evaluate all models: {elapsed_time} seconds")

# Plot training accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(train_accuracies[name], label=f'{name}')
plt.title('Training Accuracy by w/wo Dropout')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy')
plt.legend()
plt.show()

# Plot testing accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(test_accuracies[name], label=f'{name}')
plt.title('Testing Accuracy by w/wo Dropout')
plt.xlabel('Epoch')
plt.ylabel('Testing Accuracy')
plt.legend()
plt.show()

```

Training 4 layer RELU:

Epoch #10:	train accuracy 0.897	train loss 0.271	test accuracy
0.790	test loss 0.486		
Epoch #20:	train accuracy 0.962	train loss 0.130	test accuracy
0.767	test loss 0.672		
Epoch #30:	train accuracy 0.995	train loss 0.043	test accuracy
0.771	test loss 0.916		
Epoch #40:	train accuracy 1.000	train loss 0.011	test accuracy
0.768	test loss 1.259		
Epoch #50:	train accuracy 1.000	train loss 0.002	test accuracy
0.774	test loss 1.516		
Epoch #60:	train accuracy 1.000	train loss 0.001	test accuracy

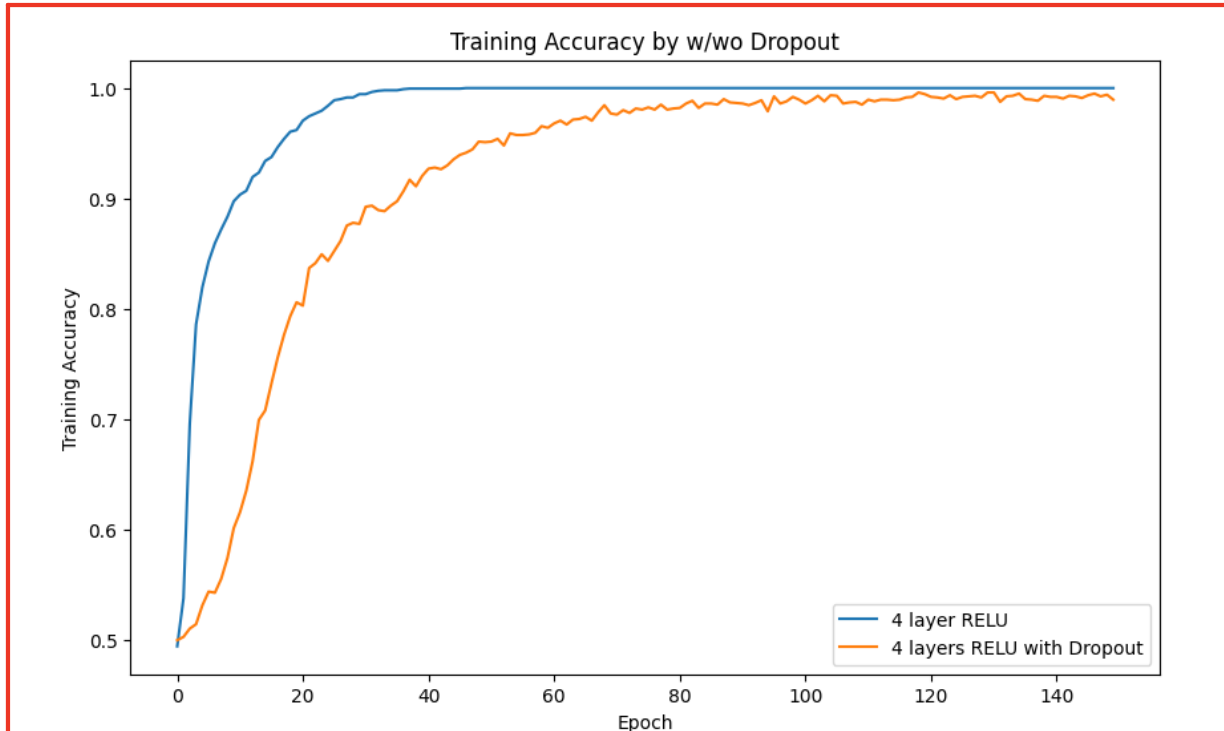
0.776	test loss 1.795		
Epoch #70:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.021		
Epoch #80:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 2.231		
Epoch #90:	train accuracy 1.000	train loss 0.000	test accuracy
0.775	test loss 2.434		
Epoch #100:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.629		
Epoch #110:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 2.810		
Epoch #120:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 3.000		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.774	test loss 3.194		
Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.772	test loss 3.370		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.773	test loss 3.541		


Training 4 layers RELU with Dropout:

Epoch #10:	train accuracy 0.602	train loss 0.671	test accuracy
0.680	test loss 0.668		
Epoch #20:	train accuracy 0.806	train loss 0.477	test accuracy
0.762	test loss 0.515		
Epoch #30:	train accuracy 0.877	train loss 0.325	test accuracy
0.782	test loss 0.481		
Epoch #40:	train accuracy 0.920	train loss 0.239	test accuracy
0.784	test loss 0.520		
Epoch #50:	train accuracy 0.951	train loss 0.155	test accuracy
0.786	test loss 0.593		
Epoch #60:	train accuracy 0.964	train loss 0.113	test accuracy
0.784	test loss 0.703		
Epoch #70:	train accuracy 0.977	train loss 0.071	test accuracy
0.780	test loss 0.840		
Epoch #80:	train accuracy 0.982	train loss 0.059	test accuracy
0.785	test loss 1.005		
Epoch #90:	train accuracy 0.987	train loss 0.043	test accuracy
0.779	test loss 1.112		
Epoch #100:	train accuracy 0.990	train loss 0.033	test accuracy
0.784	test loss 1.264		
Epoch #110:	train accuracy 0.985	train loss 0.040	test accuracy
0.784	test loss 1.324		
Epoch #120:	train accuracy 0.995	train loss 0.021	test accuracy
0.782	test loss 1.467		
Epoch #130:	train accuracy 0.996	train loss 0.018	test accuracy
0.783	test loss 1.515		
Epoch #140:	train accuracy 0.992	train loss 0.029	test accuracy

0.784 test loss 1.636
Epoch #150: train accuracy 0.990 train loss 0.024 test accuracy
0.783 test loss 1.737

Time to train and evaluate all models: 107.53877091407776 seconds





Analysis and discussion here (< 5 sentences) : The training accuracy plot demonstrates that the four-layer network without dropout achieves higher training accuracy more rapidly compared to the same network with dropout. This suggests that dropout regularizes the model by limiting its capacity to overfit to the training data. On the testing accuracy plot, both models initially show similar accuracy, but as epochs increase, the model with dropout maintains a more consistent accuracy. This consistency indicates effective regularization, which helps the model generalize better without reducing its overall performance on unseen data. Hence, while dropout seems to slow down learning, it leads to a model that is more robust to variations in the data.

3.4 Question 2.4 (5 points)

Now that you have trained a Neural Network classifier, you may find that your testing accuracy is much lower than the training accuracy. In what ways can we decrease this gap? Pick all that apply.

1. Train on a larger dataset.
2. Add more hidden units.
3. Increase the regularization strength.
4. None of the above.

Answer here :

- Train on a larger dataset.
- Increase the regularization strength.

Explanation (< 5 sentences) here: : Overfitting occurs when a model is too complex relative to the amount and noisiness of the input data, capturing spurious patterns that fail to generalize to other data sets. More data and stronger regularization both work to smooth out these issues, improving model performance on new, unseen datasets (testing accuracy) relative to the performance on the training dataset. Adding more hidden units, conversely, might increase the model's complexity and potentially lead to more overfitting, not less.

4 Part 3: Exploration (20 points)

4.1 Question 3.1 Explore (20 points)

There are other aspects to optimizing neural network performance. Explore two here, and discuss your findings. You may also try different neural architectures here, other than feedforward networks.

```
[16]: ### YOUR CODE HERE

# Learning Rate Scheduling
import torch
import torch.nn as nn
import torch.optim as optim
from torch.optim.lr_scheduler import StepLR

def experiment_steplr(model):
    loss_fn = nn.NLLLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=0.0001)
    scheduler = StepLR(optimizer, step_size=30, gamma=0.1) # Adjust lr every
    ↪ 30 epochs

    all_train_accuracy = []
    average_train_loss = []
    all_test_accuracy = []
    average_test_loss = []

    for epoch in range(150):
        train_accuracy, train_loss = train_epoch(train_loader, model, loss_fn,
    ↪ optimizer)
        all_train_accuracy.append(train_accuracy)
        average_train_loss.append(train_loss)

        test_accuracy, test_loss = eval_epoch(test_loader, model, loss_fn,
    ↪ optimizer)
        all_test_accuracy.append(test_accuracy)
        average_test_loss.append(test_loss)

        scheduler.step() # Update the learning rate

    if epoch % 10 == 9:
```

```

        print(f'Epoch #{epoch+1}: \t train accuracy {train_accuracy:.3f} \t \t
↳train loss {train_loss:.3f} \t test accuracy {test_accuracy:.3f} \t test \t
↳loss {test_loss:.3f}')
```

```

    return all_train_accuracy, all_test_accuracy

# Batch Normalization
class NN2_BatchNorm(nn.Module):
    def __init__(self, input_size, hidden_size=100):
        super(NN2_BatchNorm, self).__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.bn1 = nn.BatchNorm1d(hidden_size)
        self.fc2 = nn.Linear(hidden_size, 2)

    def forward(self, x):
        x = F.relu(self.bn1(self.fc1(x)))
        x = F.log_softmax(self.fc2(x), dim=1)
        return x

input_size = 512
hidden_size = 100
models = [NN2(input_size, hidden_size), NN2_BatchNorm(input_size, hidden_size),
↳NN2(input_size, hidden_size)]
model_names = ['2 layers', '2 layers with Batch Normalization', '2 layers with \t
↳Learning Rate Scheduling']

# Dictionaries to hold training and testing accuracies
train_accuracies = {}
test_accuracies = {}

start_time = time.time()

# Loop through each model, train, and collect accuracies
for model, name in zip(models, model_names):
    print(f'\nTraining {name}:')
    if name == '2 layers with Learning Rate Scheduling':
        train_accuracy, test_accuracy = experiment_steplr(model)
    else:
        train_accuracy, test_accuracy = experiment(model)

    train_accuracies[name] = train_accuracy
    test_accuracies[name] = test_accuracy

end_time = time.time()
elapsed_time = end_time - start_time
print(f"\nTime to train and evaluate all models: {elapsed_time} seconds")

```

```

# Plot training accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(train_accuracies[name], label=f'{name}')
plt.title('Training Accuracy by w/o Batch Normalization')
plt.xlabel('Epoch')
plt.ylabel('Training Accuracy')
plt.legend()
plt.show()

# Plot testing accuracies
plt.figure(figsize=(10, 6))
for name in model_names:
    plt.plot(test_accuracies[name], label=f'{name}')
plt.title('Testing Accuracy by w/o Batch Normalization')
plt.xlabel('Epoch')
plt.ylabel('Testing Accuracy')
plt.legend()
plt.show()

```

Training 2 layers:

Epoch #10:	train accuracy 0.897	train loss 0.292	test accuracy
0.787	test loss 0.476		
Epoch #20:	train accuracy 0.944	train loss 0.182	test accuracy
0.786	test loss 0.516		
Epoch #30:	train accuracy 0.977	train loss 0.116	test accuracy
0.776	test loss 0.571		
Epoch #40:	train accuracy 0.993	train loss 0.071	test accuracy
0.777	test loss 0.671		
Epoch #50:	train accuracy 0.998	train loss 0.042	test accuracy
0.772	test loss 0.781		
Epoch #60:	train accuracy 0.999	train loss 0.024	test accuracy
0.770	test loss 0.857		
Epoch #70:	train accuracy 1.000	train loss 0.013	test accuracy
0.770	test loss 0.958		
Epoch #80:	train accuracy 1.000	train loss 0.007	test accuracy
0.766	test loss 1.078		
Epoch #90:	train accuracy 1.000	train loss 0.004	test accuracy
0.765	test loss 1.196		
Epoch #100:	train accuracy 1.000	train loss 0.002	test accuracy
0.762	test loss 1.336		
Epoch #110:	train accuracy 1.000	train loss 0.001	test accuracy
0.765	test loss 1.434		
Epoch #120:	train accuracy 1.000	train loss 0.001	test accuracy
0.764	test loss 1.586		
Epoch #130:	train accuracy 1.000	train loss 0.000	test accuracy
0.762	test loss 1.678		

Epoch #140:	train accuracy 1.000	train loss 0.000	test accuracy
0.761	test loss 1.805		
Epoch #150:	train accuracy 1.000	train loss 0.000	test accuracy
0.762	test loss 1.926		

Training 2 layers with Batch Normalization:


Epoch #10:	train accuracy 0.907	train loss 0.271	test accuracy
0.795	test loss 0.463		
Epoch #20:	train accuracy 0.948	train loss 0.165	test accuracy
0.780	test loss 0.484		
Epoch #30:	train accuracy 0.972	train loss 0.104	test accuracy
0.767	test loss 0.526		
Epoch #40:	train accuracy 0.987	train loss 0.065	test accuracy
0.764	test loss 0.618		
Epoch #50:	train accuracy 0.990	train loss 0.048	test accuracy
0.762	test loss 0.680		
Epoch #60:	train accuracy 0.995	train loss 0.032	test accuracy
0.760	test loss 0.754		
Epoch #70:	train accuracy 0.987	train loss 0.042	test accuracy
0.764	test loss 0.772		
Epoch #80:	train accuracy 0.994	train loss 0.027	test accuracy
0.761	test loss 0.807		
Epoch #90:	train accuracy 1.000	train loss 0.016	test accuracy
0.754	test loss 0.847		
Epoch #100:	train accuracy 0.994	train loss 0.021	test accuracy
0.765	test loss 0.870		
Epoch #110:	train accuracy 1.000	train loss 0.012	test accuracy
0.761	test loss 0.912		
Epoch #120:	train accuracy 0.999	train loss 0.010	test accuracy
0.750	test loss 0.979		
Epoch #130:	train accuracy 0.997	train loss 0.015	test accuracy
0.745	test loss 1.037		
Epoch #140:	train accuracy 0.998	train loss 0.011	test accuracy
0.751	test loss 0.963		
Epoch #150:	train accuracy 1.000	train loss 0.006	test accuracy
0.757	test loss 1.008		


Training 2 layers with Learning Rate Scheduling:

Epoch #10:	train accuracy 0.892	train loss 0.302	test accuracy
0.795	test loss 0.468		
Epoch #20:	train accuracy 0.935	train loss 0.196	test accuracy
0.784	test loss 0.520		
Epoch #30:	train accuracy 0.969	train loss 0.128	test accuracy
0.786	test loss 0.576		
Epoch #40:	train accuracy 0.978	train loss 0.114	test accuracy
0.786	test loss 0.583		
Epoch #50:	train accuracy 0.979	train loss 0.109	test accuracy
0.783	test loss 0.591		

Epoch #60:	train accuracy 0.981	train loss 0.103	test accuracy
0.784	test loss 0.602		
Epoch #70:	train accuracy 0.981	train loss 0.102	test accuracy
0.783	test loss 0.603		
Epoch #80:	train accuracy 0.981	train loss 0.101	test accuracy
0.783	test loss 0.604		
Epoch #90:	train accuracy 0.982	train loss 0.101	test accuracy
0.783	test loss 0.605		
Epoch #100:	train accuracy 0.982	train loss 0.101	test accuracy
0.783	test loss 0.605		
Epoch #110:	train accuracy 0.982	train loss 0.100	test accuracy
0.782	test loss 0.605		
Epoch #120:	train accuracy 0.982	train loss 0.100	test accuracy
0.782	test loss 0.605		
Epoch #130:	train accuracy 0.982	train loss 0.100	test accuracy
0.782	test loss 0.606		
Epoch #140:	train accuracy 0.982	train loss 0.100	test accuracy
0.782	test loss 0.606		
Epoch #150:	train accuracy 0.982	train loss 0.100	test accuracy
0.782	test loss 0.606		

Time to train and evaluate all models: 144.76883625984192 seconds





Analysis and discussion here (< 15 sentences) : It is evident that all models achieve high training accuracy over time, but their testing accuracy does not improve in tandem. This indicates overfitting, especially for the model without batch normalization or learning rate scheduling, which shows a testing accuracy that peaks and then starts to decrease as the epochs increase. The model with batch normalization appears to stabilize the training slightly better than the one without it, but ultimately it too shows signs of overfitting, as seen by the gap between training and testing accuracy and the decrease in testing accuracy after a certain number of epochs. The model with learning rate scheduling shows the best testing accuracy overall, although the improvement is modest. It seems to slightly mitigate overfitting as the testing accuracy does not degrade as much as with the other models. However, the plateau in testing accuracy despite further training suggests that simply adjusting the learning rate is not enough to continue to improve the generalization of the model. In all cases, while training accuracy approaches or reaches 100%, the testing accuracy is significantly lower, highlighting the models' inability to generalize well from their training data to unseen data. This could be improved by employing techniques such as data augmentation, more sophisticated learning rate schedules, early stopping, or by exploring more complex models that may capture the nuances of the data better without overfitting.

Submission Instructions

1. Click the Save button at the top of the Jupyter Notebook.

2. Select Cell -> All Output -> Clear. This will clear all the outputs from all cells (but will keep the content of all cells).
3. Select Cell -> Run All. This will run all the cells in order, and will take several minutes.
4. Once you've rerun everything, select File -> Download as -> PDF via LaTeX (If you have trouble using "PDF via LaTeX", you can also save the webpage as pdf. Make sure all your solutions are displayed in the pdf, it's okay if the provided codes get cut off because lines are not wrapped in code cells).
5. Look at the PDF file and make sure all your solutions are there, displayed correctly. The PDF is the only thing your graders will see!
6. Submit your PDF on Gradescope.